

Michael Hodges

All files can be found here

https://github.com/Michael-Hodges/EECE5644_Machine_Learning.git
or in the appendix

Problem 1

Conduct the following model order selection exercise using 10-fold cross-validation procedure and report your procedure and results in a comprehensive, convincing, and rigorous fashion:

1. Select a Gaussian Mixture Model as the true probability density function for 2-dimensional real-valued data synthesis. This GMM will have 4 components with different mean vectors, different covariance matrices, and different probability for each Gaussian to be selected as the generator for each sample. Specify the true GMM that generates data.
2. Generate multiple data sets with independent identically distributed samples using this true GMM; these datasets will have respectively 10, 100, 1000, 10000 samples.
3. For each data set, using maximum likelihood parameter estimation principles (e.g. with the EM algorithm), within the framework of $K(=10)$ -fold cross-validation, evaluate GMMs with different model orders; specifically evaluate candidate GMMs with 1, 2, 3, 4, 5, 6 Gaussian components. Note that both model parameter estimation and validation performance measures to be used is log-likelihood of data.
4. Report your results for the experiment, indicating which of the six GMM orders get selected for each of the datasets you produced. Develop a good way to describe and summarize your experiment results in the form of tables/figures.

Our true model is as follows:

Gaussian	Mean	Covariance	Prior
1	[2, 2]	[1 0; 0 1]	0.2
2	[-2, 2]	[0.1 0; 0 2]	0.3
3	[-2, -2]	[2 0; 0 0.1]	0.1
4	[2, -2]	[1 -0.7; -0.7 1]	0.4

The distribution is shown below:

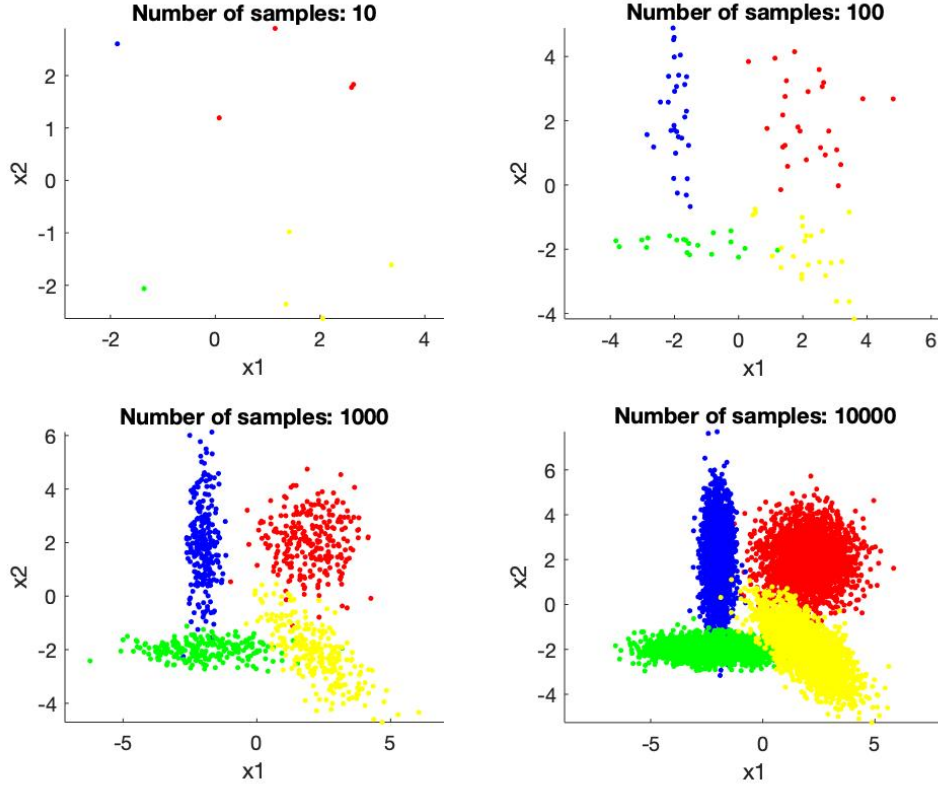


Figure 1: Original distribution for $N= 10, 100, 1000, 10000$

After performing $k=10$ cross fold validation and sweeping the number of mixture models to use from 1:6 this is our log loss distribution per fold shown in tables

-3.9638	-4.2978	-6.8071	-4.6400	-3.2910	-4.2495	-4.3731	-3.4920	-7.3494	-4.3322
-3.6247	-4.0536	-10.7713	-5.1538	-7.1021	-4.1788	-5.3541	-4.2197	-8.6037	-1.5662
-3.8892	-5.1617	-10.7713	-4.5627	-7.1024	-1.5734	-36.8864	-17.6312	-36.9400	-1.5674
-3.6245	-4.5059	-30.3148	-4.5627	-15.1959	-1.5734	-35.8748	-8.3713	-39.4311	-1.5674
-31.4103	-51.8002	-36.7608	-20.1127	-7.1024	-1.5734	-336.7273	-143.8673	-365.2600	-3.5226
-31.4103	-55.0988	-654.4114	-170.9504	-20.4518	0.8003	-186.4599	-349.1736	-575.3199	0.8003

Table 1: $N = 10$ Log loss with rows = number of gaussian models[1:6], columns = fold number [1:10]

-45.3176	-47.0058	-45.9258	-41.7257	-47.7805	-43.8224	-45.0922	-43.6180	-42.0198	-44.8303
-42.0365	-45.5029	-42.7794	-39.7459	-45.8585	-40.5367	-42.9307	-42.1857	-39.1282	-45.0525
-43.8706	-47.0594	-36.7414	-40.3071	-46.0681	-35.0121	-40.5685	-39.0395	-36.8316	-41.0783
-42.2338	-45.4390	-34.7023	-37.1690	-43.0764	-32.7697	-45.5257	-41.8924	-34.8293	-39.5333
-41.9808	-50.6697	-38.8393	-39.2670	-46.2515	-34.8718	-43.1371	-41.7510	-35.6455	-42.7909
-45.1952	-43.2287	-40.0771	-39.0436	-43.5951	-32.7287	-40.3351	-41.1926	-38.3550	-40.5447

Table 2: $N = 100$ Log loss with rows = number of gaussian models[1:6], columns = fold number [1:10]

-439.0841	-450.1371	-446.3358	-448.2369	-447.9725	-452.6688	-457.8052	-454.7543	-442.5781	-445.7355
-432.9024	-426.7419	-411.9246	-419.4427	-427.0980	-427.3741	-431.7883	-435.9189	-420.0560	-427.0808
-399.7574	-383.6188	-377.5681	-387.3891	-408.3147	-389.9040	-396.0273	-400.4866	-378.3312	-395.9905
-379.3314	-364.5988	-360.9133	-373.1415	-362.6514	-370.6030	-371.4420	-375.1668	-364.2285	-378.6182
-380.2124	-365.7982	-360.3902	-373.6248	-364.0981	-371.6876	-370.7591	-374.6762	-365.0757	-377.1173
-380.6687	-367.6231	-359.9976	-374.1620	-362.1270	-371.6685	-372.0235	-374.9097	-363.1726	-379.3920

Table 3: N = 1000 Log loss with rows = number of gaussian models[1:6], columns = fold number [1:10]

1.0e+03 *

-4.4492	-4.4438	-4.4639	-4.4081	-4.4572	-4.4560	-4.4625	-4.4154	-4.4546	-4.4613
-4.1879	-4.1647	-4.2326	-4.2342	-4.1896	-4.1636	-4.2086	-4.2095	-4.1851	-4.2288
-3.8551	-3.8175	-3.8071	-3.8415	-3.7987	-3.8481	-3.9486	-3.8723	-3.8624	-3.9101
-3.7209	-3.6626	-3.6680	-3.7086	-3.6493	-3.6613	-3.7310	-3.7159	-3.7051	-3.7123
-3.7220	-3.6620	-3.6683	-3.7089	-3.6498	-3.6620	-3.7315	-3.7158	-3.7048	-3.7122
-3.7210	-3.6624	-3.6690	-3.7088	-3.6496	-3.6617	-3.7308	-3.7157	-3.7056	-3.7121

Table 4: N = 10000 Log loss with rows = number of gaussian models[1:6], columns = fold number [1:10]

A graph summarizing the results is shown below:

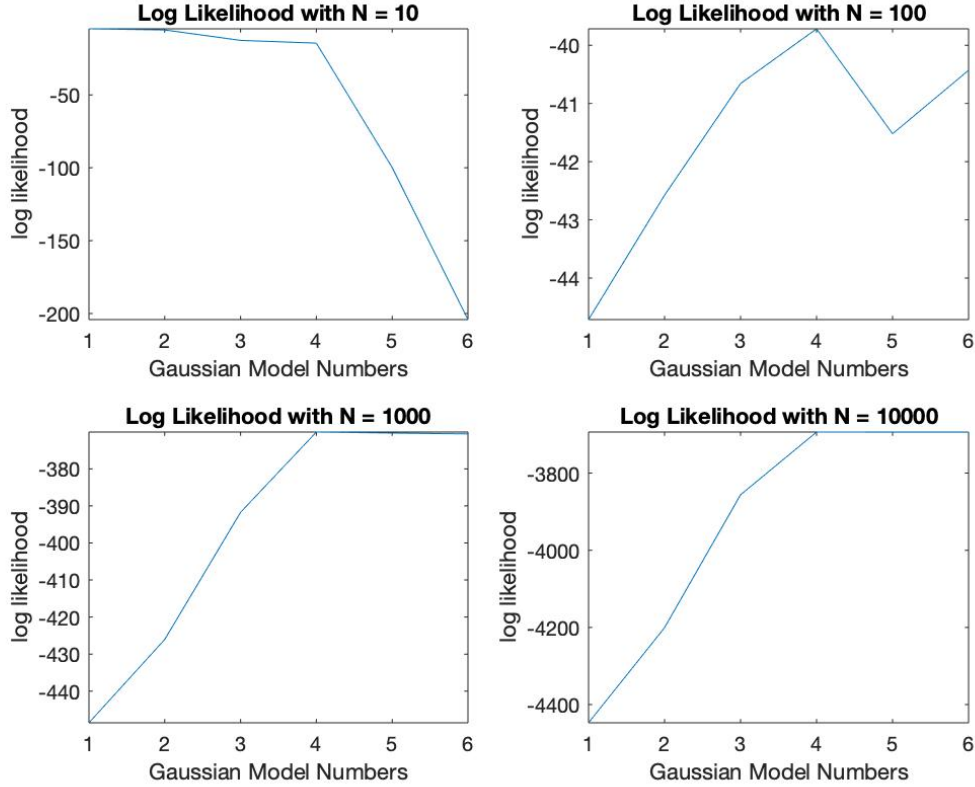


Figure 2: Log Likelihoods for N= 10, 100, 1000, 10000

As we can see after we reach a sample set of $N = 100$ our model chooses to have 4 Gaussian models in our Gaussian Mixture Model. At $N = 10$ there is not enough data to accurately find 4 models so our system chooses one that captures all of the data given the mean and covariance. Furthermore, for $N = 1,000$ and $10,000$ there appears to about be a tie between 4-6 models, but we would choose the simplest model to represent our solution so we would choose 4 in this case.

Problem 2

Conduct the following maximum likelihood discriminative classifier training exercise on data generated from two Gaussian distributed classes:

1. Generate training set with 999 2-dimensional samples from two classes with priors $q_- = 0.3$ and $q_+ = 0.7$; the class-conditional data probability distributions are two Gaussians with different mean vectors and different covariance matrices (choose the matrices to be non-diagonal with distinct eigenvalues, so your Gaussian pdfs are tilted with respect to each other and elongated in different directions by different aspect ratios). Hint: For more interesting results, make the Gaussians overlap with each other somewhat significantly, so that the minimum error probability achievable is not too small.
2. Using Fisher LDA, identify a linear classifier that minimizes the error count on the training set. This classifier will have a discriminant function of the form $w^T \text{LDA}x + b_{\text{LDA}}$ where x is the data vector, and the classifier decides in favor of Class - if the discriminant is below 0, and decides in favor of Class +, if the discriminant is at least 0.
3. Train the parameters of a logistic function $y(x) = 1/(1+e^{-(w^T x + b)})$ using the maximum likelihood estimation principle to optimize the parameters w and b with the training set, such that the function is trained to act as a surrogate for the posterior probability of Class + given x . In particular, your model assumes that $y(x) = P(\text{Label} = +|x)$; consequently, $1 - y(x) = P(\text{Label} = -|x)$. Hint: Once you specify the optimization objective to train this logistic-linear-model for class posterior, you can solve the optimization problem using any suitable numerical optimization procedure, such as gradient ascent that you implement from scratch, or using a derivative free numerical optimization procedure like the Nelder-Mead Simplex Reflection Algorithm (e.g. in Matlab, `fminsearch`). Make a choice, implement correctly, perhaps consider using the LDA solution you developed earlier to provide an initial estimate for the model parameters.
4. Report visual and numerical results that compare the following three classifiers (e.g. data scatter plots with color/shape indicators of true/decided labels), including the error counts each classifier achieve on the training set: MAP-classifier that makes use of the true data distributions and class priors, which achieves minimum probability error by design; LDA classifier you designed earlier; logistic-linear classifier you designed next.

$$\begin{aligned}
\hat{\theta}_{ML} &= \operatorname{argmax}_{\Theta} \sum_{i=1}^N \ln P(x_i, l_i | \Theta) \\
&= \operatorname{argmax}_{\Theta} \sum_{i=1}^N \ln [(P(l_i | x_i, \Theta) P(x_i | \Theta))] \\
&= \operatorname{argmax}_{\Theta} \sum_{i=1}^N \ln [(P(l_i | x_i, \Theta) P(x_i | \Theta))] \\
&= \operatorname{argmax}_{\Theta} \sum_{i=1}^N \ln (P(l_i | x_i, \Theta) + \sum_{i=1}^N \ln P(x_i | \Theta)) \\
&= \operatorname{argmax}_{\Theta} \sum_{i=1}^N \ln y(x)^{q_i} (1 - y(x))^{1-q_i} \\
&= \operatorname{argmax}_{\Theta} \sum_{i=1}^N q_i \ln y(x) + (1 - q_i) \ln(1 - y(x)) \\
&= \operatorname{argmax}_{\mathbf{w}, b} \sum_{i=1}^N q_i \ln \left(\frac{1}{1 + e^{(\mathbf{w}^T \mathbf{x} + b)}} \right) + (1 - q_i) \ln \left(1 - \frac{1}{1 + e^{(\mathbf{w}^T \mathbf{x} + b)}} \right)
\end{aligned}$$

Where q_i is the prior probability for the class x belongs to. Note: we essentially use a bernoulli to solve this.

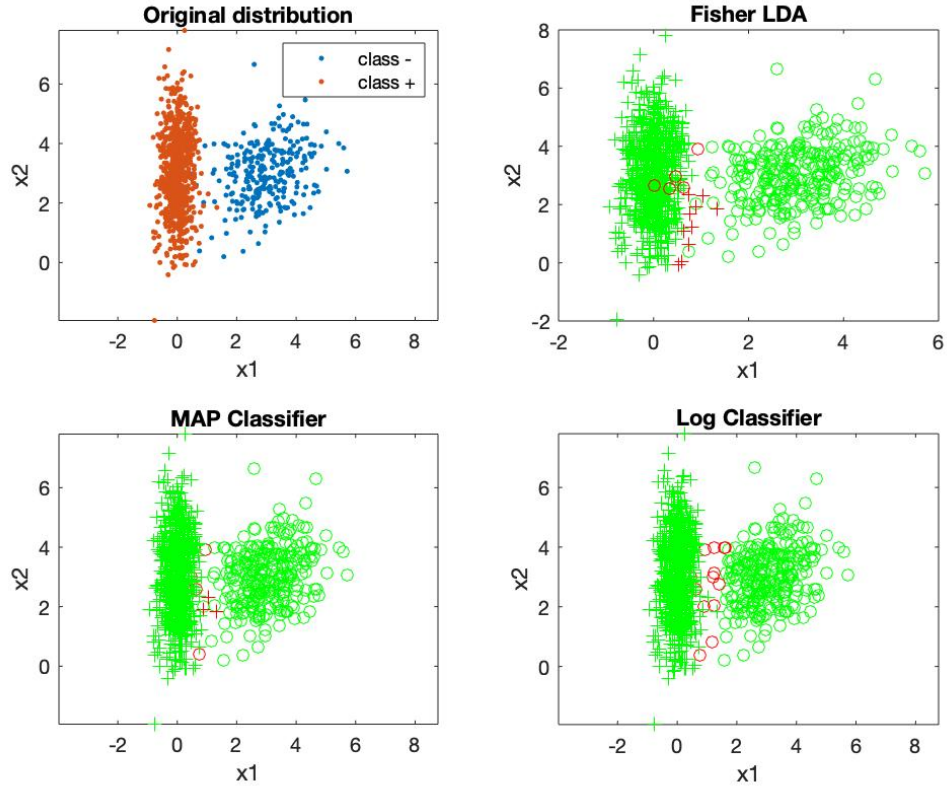


Figure 3: True class labels and classification. 0 class -, + Class +, Red incorrect, green correct

FischerLDA	MAP	Log
1.5%	0.9%	1.5%

Table 5: Percent mislabeled per model

FischerLDA	MAP	Log
15	9	15

Table 6: Number mislabeled per model

Appendix

1. Problem 1

```
1 clear all; close all; clc
2 % Gaussian Mixture Model Parameters
3 % Means:
4 mu(:,1) = [2, 2];
5 mu(:,2) = [-2, 2];
6 mu(:,3) = [-2, -2];
7 mu(:,4) = [2, -2];
8
9 % Covariances:
10 Sigma(:, :, 1) = [1 0; 0 1];
11 Sigma(:, :, 2) = [0.1 0; 0 2];
12 Sigma(:, :, 3) = [2 0; 0 0.1];
13 Sigma(:, :, 4) = [1 -0.7; -0.7 1];
14
15 % Priors:
16 classPriors = [0.2 0.3 0.1 0.4]; % True Priors
17 classPriors1 = [0.25 0.25 0.25 0.251];
18 assert(max(cumsum(classPriors)) == 1, 'Priors do not equal 1');
19 thr = [0, cumsum(classPriors1)];
20
21 disp("True Means");
22 disp(mu);
23 disp("True Covariances:");
24 disp(Sigma);
25 disp("True Priors:");
26 disp(classPriors);
27 k = 10; % Perform k fold validation
28
29 %Generate Data:
30 % figure(1), clf, colorList = 'rbgy';
31 for i = 1:4
32     N = 10^i; % Number of samples to generate for each set
33     (10,100,1000,10000)
34     u = rand(1,N); L = zeros(1,N); x = zeros(2,N);
35     for l = 1:4
36         indices = find(thr(l) <= u & u < thr(l+1)); % fixed using
37             classPriors1 adding a small term to last prior if u
38             happens to be precisely 1, that sample will get
39             omitted - needs to be fixed
40         L(1, indices) = l * ones(1, length(indices));
41         x(:, indices) = mvnrnd(mu(:, l), Sigma(:, :, l), length(
42             indices));
43     end
44     % subplot(2,2,i);
45     % plot(x(1, indices), x(2, indices), '.', 'MarkerFaceColor',
```



```

colorList(1)); axis equal, hold on,
40 %         axis([-10 10 -10 10]);
41 end
42 %     hold off;
43     assert( isempty( find(L==0, 1) ), 'some values unclassified' );
44     x(3,:) = L;
45     disp("Number Samples N = " + 10^i);
46     disp("Class1: " + num2str(length( find(L==1))));
47     disp("Class2: " + num2str(length( find(L==2))));
48     disp("Class3: " + num2str(length( find(L==3))));
49     disp("Class4: " + num2str(length( find(L==4))));
50
51
52     % Loop through number of models to try 1-6
53     regWeight = 1e-10; % regularization parameter for
        covariance estimates
54     for l = 1:6 % number of models to loop through
55         disp("performing: " + k + "-fold cross validation N=" +
            10^i + ", Models =" + l);
56         for q = 1:k % k fold cross validation
57             disp(q+"th fold");
58             dVal = x(:,1+(N/k)*(q-1):N/k*q);
59 %             disp("All " + num2str(x));
60 %             disp("Validation: " + num2str(dVal));
61             dTrain = x;
62             dTrain(:,1+(N/k)*(q-1):N/k*q) = [];
63 %             disp("Train: " + num2str(dTrain));
64             % Set up initial GMM for l number of Gaussians
65             EMdTrain = dTrain; % take train and set equal to
                new set to work with
66             shuffledIndices = randperm(length(EMdTrain(1,:)));
                % Pick l random samples to be initial mean
67             EM_mu = EMdTrain(1:2,shuffledIndices(1:l));
                EM_mu_new = EM_mu; % starting means
68 %             [~,EMdTrain(3,:)] = min(pdist2(EM_mu',EMdTrain
(1:2,:)'),[],1); % Set labels to nearest centroid
69 %             EM_Sigma = zeros(2,2,l); EM_Sigma_new = EM_Sigma;
70 %             EM_Prior = zeros(1,l); EM_Prior_new = EM_Prior;
71             % initialize Priors
72 %             EM_Prior = ones(1,l)./l;
73 %             for b = 1:l
74 %                 EM_Prior(b) = length( find(EMdTrain(3,:)==b) ) /
length(EMdTrain(3,:));
75 %             end
76             % Initialize Covariances
77 %             for b = 1:l

```

```

78 %             EM_Sigma(:, :, b) = cov(EMdTrain(1:2, find(
EMdTrain(3, :) == b))) + regWeight * eye(2, 2); % Generate
covariances with given labels
79 %             end
80 %             epsilon = 1e-4; % convergence factor
81 %             Converged = 0;
82 %             t = 0;
83 % Try two initializations
84 [EMlog1, EM_Prior1, EM_mu1, EM_Sigma1] =
EMforGMM_Edited(length(EMdTrain(3, :)), EMdTrain
(1:2, :), 1);
85 [EMlog2, EM_Prior2, EM_mu2, EM_Sigma2] =
EMforGMM_Edited(length(EMdTrain(3, :)), EMdTrain
(1:2, :), 1);
86 if EMlog1 > EMlog2
87     EM_Prior = EM_Prior1;
88     EM_mu = EM_mu1;
89     EM_Sigma = EM_Sigma1;
90 else
91     EM_Prior = EM_Prior2;
92     EM_mu = EM_mu2;
93     EM_Sigma = EM_Sigma2;
94 end
95 loglikelihood(l, q, i) = sum(log(evalGMM(dVal(1:2, :),
EM_Prior, EM_mu, EM_Sigma))));
96 end
97 disp(loglikelihood(:, :, i));
98 %Store error here for each model
99
100 end
101 figure(2); hold on;
102 subplot(2, 2, i);
103 means = mean(loglikelihood(:, :, i), 2);
104 plot(1:6, means');
105 xlabel("Gaussian Model Numbers");
106 ylabel("log likelihood");
107 title("Log Likelihood with N = " + 10^i);
108 axis([1 6 -inf inf]);
109 rangex1 = [min(x(1, :)), max(x(1, :))];
110 rangex2 = [min(x(2, :)), max(x(2, :))];
111 [x1Grid, x2Grid, zGMM] = contourGMM(EM_Prior, EM_mu, EM_Sigma,
rangex1, rangex2);
112 figure(3);
113 subplot(2, 2, i); hold on;
114 plot(x(1, find(L==1)), x(2, find(L==1)), 'r');
115 plot(x(1, find(L==2)), x(2, find(L==2)), 'b');

```

```

116     plot(x(1,find(L==3)),x(2,find(L==3)),'.g');
117     plot(x(1,find(L==4)),x(2,find(L==4)),'.y');
118 %     contour(x1Grid,x2Grid,zGMM);
119     title("Number of samples: " + 10^i);
120     xlabel("x1");
121     ylabel("x2");
122     axis('equal');
123 end
124
125 function gmm = evalGMM(x,alpha,mu,Sigma)
126 gmm = zeros(1,size(x,2));
127 for m = 1:length(alpha) % evaluate the GMM on the grid
128     gmm = gmm + alpha(m)*evalGaussian(x,mu(:,m),Sigma(:, :,m));
129 end
130 end
131
132 function g = evalGaussian(x,mu,Sigma)
133 % Evaluates the Gaussian pdf N(mu,Sigma) at each column of X
134 [n,N] = size(x);
135 invSigma = inv(Sigma);
136 C = (2*pi)^(-n/2) * det(invSigma)^(1/2);
137 E = -0.5*sum((x-repmat(mu,1,N)).*(invSigma*(x-repmat(mu,1,N))),1);
138 g = C*exp(E);
139 end
140
141 function [x1Grid,x2Grid,zGMM] = contourGMM(alpha,mu,Sigma,
    rangex1,rangex2)
142 x1Grid = linspace(floor(rangex1(1)),ceil(rangex1(2)),101);
143 x2Grid = linspace(floor(rangex2(1)),ceil(rangex2(2)),91);
144 [h,v] = meshgrid(x1Grid,x2Grid);
145 GMM = evalGMM([h(:)';v(:)'],alpha,mu,Sigma);
146 zGMM = reshape(GMM,91,101);
147 end

```

2. Problem 1 EM

```

1 function [logLikelihood, alpha,mu,Sigma] = EMforGMM_Edited(N,
    data, parameters)
2 % Generates N samples from a specified GMM,
3 % then uses EM algorithm to estimate the parameters
4 % of a GMM that has the same number of components
5 % as the true GMM that generates the samples.
6
7 % close all,
8 delta = 1e-2; % tolerance for EM stopping criterion
9 regWeight = 5e-3; % regularization parameter for covariance

```

```

    estimates
10
11 % Generate samples from a 3-component GMM
12 % alpha_true = [0.2,0.3,0.5];
13 % mu_true = [-10 0 10;0 0 0];
14 % Sigma_true(:, :, 1) = [3 1;1 20];
15 % Sigma_true(:, :, 2) = [7 1;1 2];
16 % Sigma_true(:, :, 3) = [4 1;1 16];
17 % x = randGMM(N, alpha_true, mu_true, Sigma_true);
18 % [d,M] = size(mu_true); % determine dimensionality of samples
    and number of GMM components
19 d=2;
20 M = parameters;
21 x = data;
22 % Initialize the GMM to randomly selected samples
23 alpha = ones(1,M)/M;
24 shuffledIndices = randperm(N);
25 mu = x(:, shuffledIndices(1:M)); % pick M random samples as
    initial mean estimates
26 [~, assignedCentroidLabels] = min(pdist2(mu', x'), [], 1); % assign
    each sample to the nearest mean
27 for m = 1:M % use sample covariances of initial assignments as
    initial covariance estimates
28     Sigma(:, :, m) = cov(x(:, find(assignedCentroidLabels==m)))' +
        regWeight*eye(d,d);
29 end
30 t = 0; %displayProgress(t,x,alpha,mu,Sigma);
31
32 Converged = 0; % Not converged at the beginning
33 while ~Converged
34     for l = 1:M % multiply prior with PDF evaluated at certain
        point
35         temp(1,:) = repmat(alpha(l), 1, N).*evalGaussian(x, mu(:, l),
            Sigma(:, :, l));
36     end
37     plgivenx = temp./sum(temp, 1); % Probability of each given
        data
38     alphaNew = mean(plgivenx, 2); % New Priors
39     w = plgivenx./repmat(sum(plgivenx, 2), 1, N); % Probability of
        given data normalized
40     muNew = x*w'; % update means mutliplying data by prior
        distributions
41     for l = 1:M
42         v = x-repmat(muNew(:, l), 1, N);
43         u = repmat(w(l, :), d, 1).*v;
44         SigmaNew(:, :, l) = u*v' + regWeight*eye(d,d); % adding a

```

```

                                small regularization term
45     end
46     Dalpha = sum(sum(abs(alphaNew-alpha)));
47     Dmu = sum(sum(abs(muNew-mu)));
48     DSigma = sum(sum(sum(abs(abs(SigmaNew-Sigma)))));
49     alpha = alphaNew; mu = muNew; Sigma = SigmaNew;
50     Converged = ((Dalpha+Dmu+DSigma)<delta); % Check if
        converged
51     t = t+1;
52     % displayProgress(t,x,alpha,mu,Sigma);
53 end
54 logLikelihood = sum(log(evalGMM(x,alpha,mu,Sigma)));
55 %keyboard,
56
57 %%%
58 function displayProgress(t,x,alpha,mu,Sigma)
59 figure(1),
60 if size(x,1)==2
61     subplot(1,2,1), cla,
62     plot(x(1,:),x(2:,:), 'b. ');
63     xlabel('x_1'), ylabel('x_2'), title('Data and Estimated GMM
        Contours'),
64     axis equal, hold on;
65     rangex1 = [min(x(1,:)),max(x(1,:))];
66     rangex2 = [min(x(2,:)),max(x(2,:))];
67     [x1Grid,x2Grid,zGMM] = contourGMM(alpha,mu,Sigma,rangex1,
        rangex2);
68     contour(x1Grid,x2Grid,zGMM); axis equal,
69     subplot(1,2,2),
70 end
71 logLikelihood = sum(log(evalGMM(x,alpha,mu,Sigma)));
72 plot(t,logLikelihood, 'b. '); hold on,
73 xlabel('Iteration Index'), ylabel('Log-Likelihood of Data'),
74 drawnow; pause(0.1),
75
76 %%%
77 function x = randGMM(N,alpha,mu,Sigma)
78 d = size(mu,1); % dimensionality of samples
79 cum_alpha = [0,cumsum(alpha)];
80 u = rand(1,N); x = zeros(d,N); labels = zeros(1,N);
81 for m = 1:length(alpha)
82     ind = find(cum_alpha(m)<u & u<=cum_alpha(m+1));
83     x(:,ind) = randGaussian(length(ind),mu(:,m),Sigma(:, :, m));
84 end
85
86 %%%

```

```

87 function x = randGaussian(N,mu,Sigma)
88 % Generates N samples from a Gaussian pdf with mean mu
    covariance Sigma
89 n = length(mu);
90 z = randn(n,N);
91 A = Sigma^(1/2);
92 x = A*z + repmat(mu,1,N);
93
94 %%%
95 function [x1Grid,x2Grid,zGMM] = contourGMM(alpha,mu,Sigma,
    rangex1,rangex2)
96 x1Grid = linspace(floor(rangex1(1)),ceil(rangex1(2)),101);
97 x2Grid = linspace(floor(rangex2(1)),ceil(rangex2(2)),91);
98 [h,v] = meshgrid(x1Grid,x2Grid);
99 GMM = evalGMM([h(:)';v(:)'],alpha,mu,Sigma);
100 zGMM = reshape(GMM,91,101);
101 %figure(1), contour(horizontalGrid,verticalGrid,
    discriminantScoreGrid,[minDSGV
        *[0.9,0.6,0.3],0,[0.3,0.6,0.9]*maxDSGV]); % plot equilevel
        contours of the discriminant function
102
103 %%%
104 function gmm = evalGMM(x,alpha,mu,Sigma)
105 gmm = zeros(1,size(x,2));
106 for m = 1:length(alpha) % evaluate the GMM on the grid
107     gmm = gmm + alpha(m)*evalGaussian(x,mu(:,m),Sigma(:, :, m));
108 end
109
110 %%%
111 function g = evalGaussian(x,mu,Sigma)
112 % Evaluates the Gaussian pdf N(mu,Sigma) at each column of X
113 [n,N] = size(x);
114 invSigma = inv(Sigma);
115 C = (2*pi)^(-n/2) * det(invSigma)^(1/2);
116 E = -0.5*sum((x-repmat(mu,1,N)).*(invSigma*(x-repmat(mu,1,N))),
    ,1);
117 g = C*exp(E);

```

3. Problem 2

```

1 clear all; close all; clc
2 % Gaussian Mixture Model Parameters
3 % Means:
4 mu(:,1) = [3, 3]; % Q-
5 mu(:,2) = [0, 3]; % Q+
6
7 % Covariances:

```

```

8 Sigma(:, :, 1) = [1 .4; .4 1]; % Q-
9 Sigma(:, :, 2) = [0.1 0; 0 2]; % Q+
10
11 % Priors:
12 classPriors = [0.3 0.7]; % True Priors
13 classPriors1 = [0.3 0.7];
14 assert(max(cumsum(classPriors)) == 1, 'Priors do not equal 1');
15 thr = [0, cumsum(classPriors1)];
16
17 N = 999; % Number of Samples to take
18 % Generate samples from each class
19 figure(1), clf, colorList = 'rbg';
20 subplot(2,2,1);
21 u = rand(1,N);
22 for l = 1:2
23     indices = find(thr(l) <= u & u < thr(l+1)); % fixed using
        classPriors1 adding a small term to last prior if u
        happens to be precisely 1, that sample will get omitted
        - needs to be fixed
24     L(1, indices) = l * ones(1, length(indices));
25     x(:, indices) = mvnrnd(mu(:, l), Sigma(:, :, l), length(indices))
        ');
26     plot(x(1, indices), x(2, indices), '.', 'MarkerFaceColor',
        colorList(l)); axis equal, hold on,
27     axis('equal');
28 end
29 legend('class -', 'class +');
30 title('Original distribution');
31 xlabel('x1'), ylabel('x2');
32 assert isempty(find(L==0, 1)), 'some values unclassified';
33 x(3, :) = L;
34 disp("Number Samples N = " + N);
35 disp("Class1: " + num2str(length(find(L==1))));
36 disp("Class2: " + num2str(length(find(L==2))));
37
38 mu1hat = mean(x(find(L==1))); S1hat = cov(x(find(L==1)));
39 mu2hat = mean(x(find(L==2))); S2hat = cov(x(find(L==2)));
40
41 % Sb = (mu1hat - mu2hat) * (mu1hat - mu2hat)';
42 % Sw = S1hat + S2hat;
43 Sb = (mu(:, 1) - mu(:, 2)) * (mu(:, 1) - mu(:, 2))';
44 Sw = Sigma(:, :, 1) + Sigma(:, :, 2);
45
46 [V, D] = eig(inv(Sw) * Sb);
47 [~, ind] = sort(diag(D), 'descend');
48 wLDA = V(:, ind(1)); % Fisher LDA projection vector

```

```

49
50 disp(wLDA)
51 yLDA = wLDA'*x(1:2,:); % All data projected on to the line
    spanned by wLDA
52 wLDA = sign(mean(yLDA(find(L==2)))-mean(yLDA(find(L==1))))*wLDA
    ; % ensures class1 falls on the + side of the axis
53 yLDA = sign(mean(yLDA(find(L==2)))-mean(yLDA(find(L==1))))*yLDA
    ; % flip yLDA accordingly
54 disp(wLDA)
55 % figure(2), clf,
56 % subplot(2,2,2);
57 % plot(yLDA(find(L==1)),zeros(1,length(find(L==1))),'o'), hold
    on,
58 % plot(yLDA(find(L==2)),zeros(1,length(find(L==2))),'+'), axis
    equal,
59 % legend('Class -','Class +'),
60 % title('LDA projection of data and their true labels'),
61 % xlabel('x_1'), ylabel('x_2'),
62 tau = 0;
63 decisionLDA = (yLDA >= -classPriors(1)/classPriors(2)); %
    classPriors(2)/classPriors(1)
64 decisionLDA = decisionLDA+1;
65
66 %Note: wLDA is w and classPriors(1)/classPriors(2) is b
67
68 ind00LDA = find(L==1 & decisionLDA==1);
69 ind01LDA = find(L==1 & decisionLDA==2);
70 ind10LDA = find(L==2 & decisionLDA==1);
71 ind11LDA = find(L==2 & decisionLDA==2);
72
73 disp((length(ind01LDA)+length(ind10LDA))/N);
74 % figure(3), clf,
75 % plot(yLDA(ind00),zeros(1,length(yLDA(ind00))),'og'), hold on,
76 % plot(yLDA(ind01),zeros(1,length(yLDA(ind01))),'or');
77 % plot(yLDA(ind11),zeros(1,length(yLDA(ind11))),'+g');
78 % plot(yLDA(ind10),zeros(1,length(yLDA(ind10))),'+r');
79 % legend('class - correct', 'class 0- incorrect', 'class +
    correct', 'class + incorrect')
80 % figure(4), clf,
81 subplot(2,2,2);
82 plot(x(1,ind00LDA),x(2,ind00LDA),'og'), hold on,
83 plot(x(1,ind11LDA),x(2,ind11LDA),'+g');
84 plot(x(1,ind01LDA),x(2,ind01LDA),'or');
85 plot(x(1,ind10LDA),x(2,ind10LDA),'+r');
86 title('Fisher LDA');
87 xlabel('x1'),ylabel('x2');

```



```

88 %hleg = legend('class - correct', 'class + correct', 'class -
    incorrect', 'class + incorrect');
89 %set(hleg,'fontsize',14)
90
91
92 %MAP
93
94 lambda = [0 1;1 0]; % loss values
95 gamma = (lambda(2,1)-lambda(1,1))/(lambda(1,2)-lambda(2,2)) *
    classPriors(1)/classPriors(2); %threshold
96 discriminantScore = log(mvnpdf(x(1:2,:),mu(:,2)',Sigma(:,:,2))
    )-log(mvnpdf(x(1:2,:),mu(:,1)',Sigma(:,:,1))));% - log(gamma
    );
97 decision = (discriminantScore >= log(gamma));
98 decision = decision+1;
99
100 ind00MAP = find(decision'==1 & L==1); %p00 = length(ind00)/Nc
    (1); % probability of true negative
101 ind10MAP = find(decision'==2 & L==1); %p10 = length(ind10)/Nc
    (1); % probability of false positive
102 ind01MAP = find(decision'==1 & L==2); %p01 = length(ind01)/Nc
    (2); % probability of false negative
103 ind11MAP = find(decision'==2 & L==2); %p11 = length(ind11)/Nc
    (2); % probability of true positive
104 %p(error) = [p10,p01]*Nc'/N; % probability of error,
    empirically estimated
105
106 % figure(5), % class 0 circle, class 1 +, correct green,
    incorrect red
107 subplot(2,2,3);
108 plot(x(1,ind00MAP),x(2,ind00MAP),'og'); hold on,
109 plot(x(1,ind10MAP),x(2,ind10MAP),'or'); hold on,
110 plot(x(1,ind01MAP),x(2,ind01MAP),'+r'); hold on,
111 plot(x(1,ind11MAP),x(2,ind11MAP),'+g'); hold on,
112 title('MAP Classifier');
113 xlabel('x1'); ylabel('x2');
114 axis equal,
115
116 disp((length(ind01MAP)+length(ind10MAP))/N);
117
118 % Logarithmic classifier
119 % fun = @(w)for 1+exp(w(1)*x(1:2,:)+w(2,1));
120 x(3,:) = 1;
121 w0 = [1 1 1];
122 %fun = @(w,x)sum(-classPriors(1)*log(1./(1+exp(w(1:2)*x(1:2,:)+
    w(3))))-classPriors(2)*log(1-1./(1+exp(w(1:2)*x(1:2,:)+w(3))

```

```

    ));
123 fun = @(w,x,L)sum( (-classPriors(L).*log(1./(1+exp(w*x)))) -
    ((1-classPriors(L)).*log(1-(1./(1+exp(w*x))))));
124 %fun = @(w,x)-sum(exp(w(1:2)*x(1:2,find(L==1))-w(3)))+sum(exp(w
    (1:2)*x(1:2,find(L==2))-w(3)));
125 f = @(w)fun(w,x,L);
126 w_l = fminsearch(f,w0);
127
128 logDscore = 1./(1+exp(w_l*x));
129 logDecision = (logDscore>=0.5);
130 logDecision = logDecision+1;
131
132 ind00log = find(logDecision==1 & L==1); %p00 = length(ind00)/Nc
    (1); % probability of true negative
133 ind10log = find(logDecision==2 & L==1); %p10 = length(ind10)/Nc
    (1); % probability of false positive
134 ind01log = find(logDecision==1 & L==2); %p01 = length(ind01)/Nc
    (2); % probability of false negative
135 ind11log = find(logDecision==2 & L==2); %p11 = length(ind11)/Nc
    (2); % probability of true positive
136
137 subplot(2,2,4);
138 plot(x(1,ind00log),x(2,ind00log),'og'); hold on,
139 plot(x(1,ind10log),x(2,ind10log),'or'); hold on,
140 plot(x(1,ind01log),x(2,ind01log),'+r'); hold on,
141 plot(x(1,ind11log),x(2,ind11log),'+g'); hold on,
142 title('Log Classifier');
143 xlabel('x1'); ylabel('x2');
144 axis equal,

```