

Michael Hodges

All files can be found here

https://github.com/Michael-Hodges/EECE5644_Machine_Learning.git
or in the appendix

Problem 1

Using the K-Means clustering algorithm with minimum Euclidean-distance-based assignments of samples to cluster centroids, segment the two attached color images into $K \in 2, 3, 4, 5$ segments. As the feature vector for each pixel use a 5-dimensional feature vector consisting of normalized vertical and horizontal coordinates of the pixel relative to the top-left corner of the image, as well as normalized red, green, and blue values of the image color at that pixel. Normalize value to make best use of the range of gray values at your disposal for visualization.

For each $K \in 2, 3, 4, 5$, let the algorithm assign labels to each pixel; specifically, label $l_{r,c} \in 1, \dots, K$ to the pixel located at row r and column c . Present your clustering results in the form of an image of these label values. Make sure you improve this segmentation outcome visualization by using a contrast enhancement method; for instance, assign a unique color value to each label and make your label image colored, or assign visually distinct grayscale value levels to each label value to make best use of the range of gray values at your disposal for visualization.

Repeat this segmentation exercise using GMM-based clustering. For each specific K , use the EM algorithm to fit a GMM with K components, and then use that GMM to do MAP-classification style cluster label assignments to pixels. Display results similarly for this alternative clustering method. Briefly comment on the reasons of any differences, if any.



Figure 1: Original

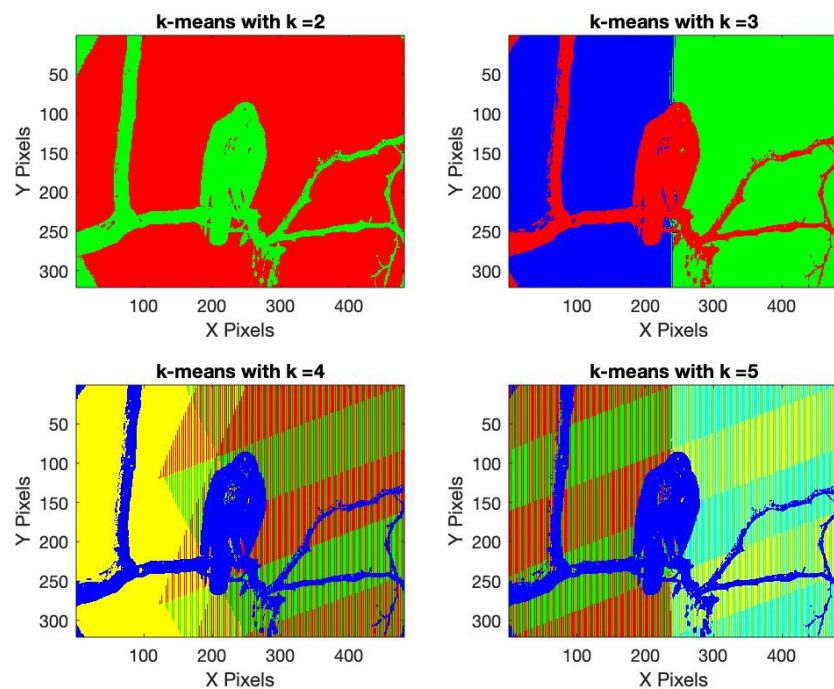


Figure 2: K-Means




Figure 3: GMM



Figure 4: Original

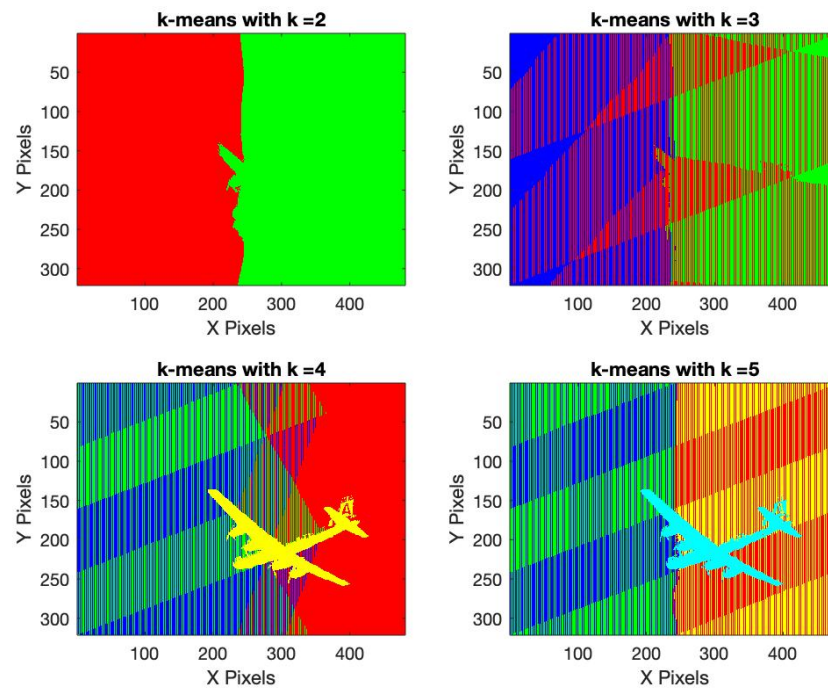
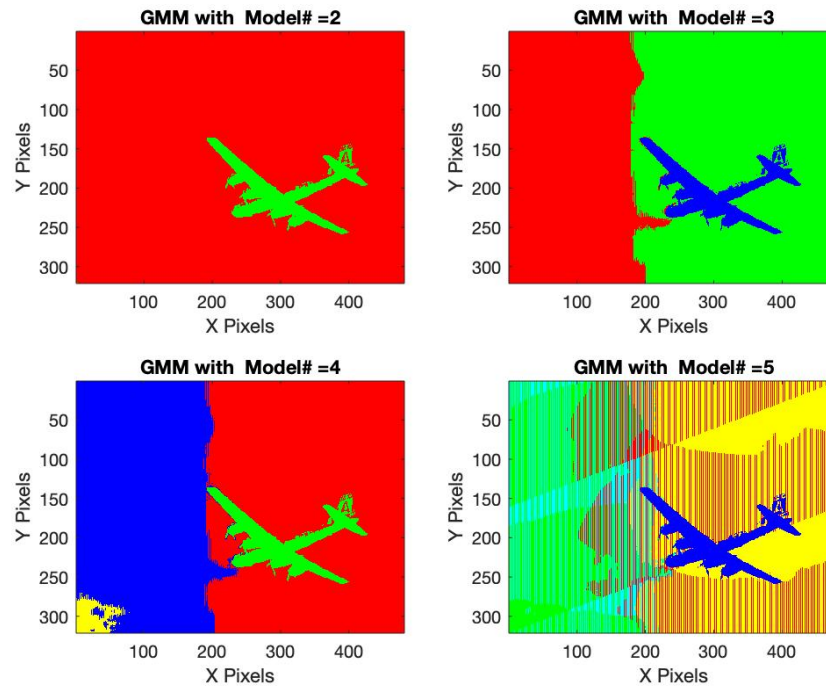


Figure 5: K-Means



As we can see the GMM generally performs better. The main reasoning is that K-Means gives a linear boundary, and based on the features we use we might not be able accurately draw a linear boundary. As we can see once we start grouping with a certain number of clusters the algorithm begins to struggle with the background of the image. Overall, we are able to pick out the main target (e.g. bird on stick, and plane) using these clustering algorithms.

Problem 2

In this exercise, you will train two support vector machine (SVM) classifiers and assess/-compare their test performances. These SVMs will respectively have linear and spherically-symmetric Gaussian (shaped radial basis function) kernels. We will refer to them as Linear-SVM and Gaussian-SVM. The data vectors are two-dimensional real-valued. The data distributions for the two classes are as follows: (1) data from class -1 are drawn from a Gaussian with zero-mean and identity- covariance-matrix; (2) data from class +1 are generated using a two-step procedure: a radius value is drawn from a uniform distribution over the interval $[2, 3]$ and an angle value (in radians) is drawn from a uniform distribution over the interval $[-\pi, \pi]$; these radius and angle values are converted to Cartesian coordinates using the Polar-to-Cartesian coordinate transformation rule.

1. Generate a training set with 1000 independent samples from these two class distributions with priors $q_- = 0.35$ and $q_+ = 0.65$; note that this does not mean 350 samples from one class and 650 from the other - the class label needs to be randomly selected for each sample, in accordance with this prior. Visualize your training data.
2. Using 10-fold cross-validation, and minimum probability of error as the objective, select the hyper parameters for both Linear-SVM and Gaussian-SVM. For both classifiers, the constraint violation term weight (usually denoted by C ; sometimes called the overlap penalty weight; referred to as the box constraint parameter in Matlab's `fitsvm`) must be optimized. For the Gaussian kernel, the scale parameter (usually denoted by σ , corresponds to the standard deviation, if this Gaussian was a probability distribution) needs to be optimized. Visualize your cross-validation process in search of optimal hyperparameter values. Report the smallest probability of error estimate you get from cross-validation.
3. Using the best hyperparameters you identified, train your Linear-SVM and Gaussian-SVM using all of the training dataset. Visualize classification results on training data, count the erroneously classified samples and report the training dataset probability of error estimate.
4. Generate 1000 independent test samples from the same class distributions with the same priors as in the training dataset. Apply the Linear-SVM and Gaussian-SVM classifiers to the test data samples. Visualize the performance of your classifiers on the test dataset and report your test probability of error estimate.

In our test we will sweep over 13 values of C , the overlap penalty weight, for the linear SVM. For the gaussian SVM we will sweep over 23 values for C , the overlap penalty, and σ , the standard deviation.




Figure 7: Train Data Original Distribution

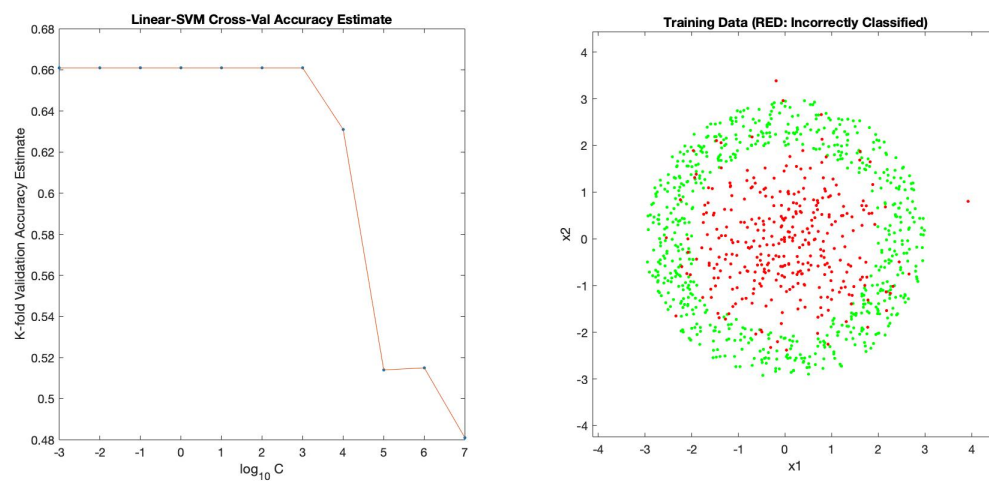


Figure 8: Train data with Linear SVM with best parameters

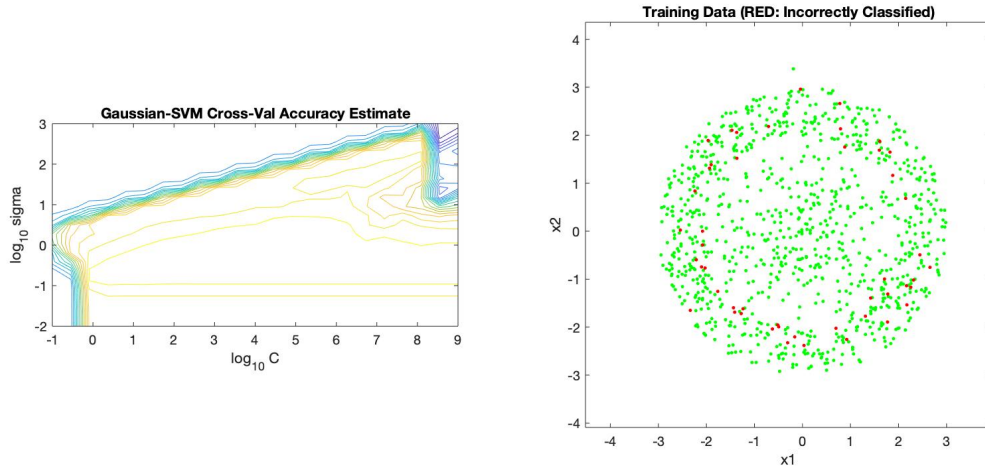


Figure 9: Train data Gaussain SVM with best parameters

From the training data we receive the following errors: Smallest Probability of error from cross validation

Linear Train Error: 33.50%

Gaussian Train Error: 4.3%

Probability of error for all data

Linear Train Error: 33.90%

Gaussian Train Error: 4.8%

This makes sense as the linear SVM shouldn't work since the data cannot be linearly seperated. Thus to reduce error we classify all of the classes as one and hence choose the class with the greatest prior. The results show what we would expect with around 35% error which matches the prior for the less likely class.




Figure 10: Test Data Original

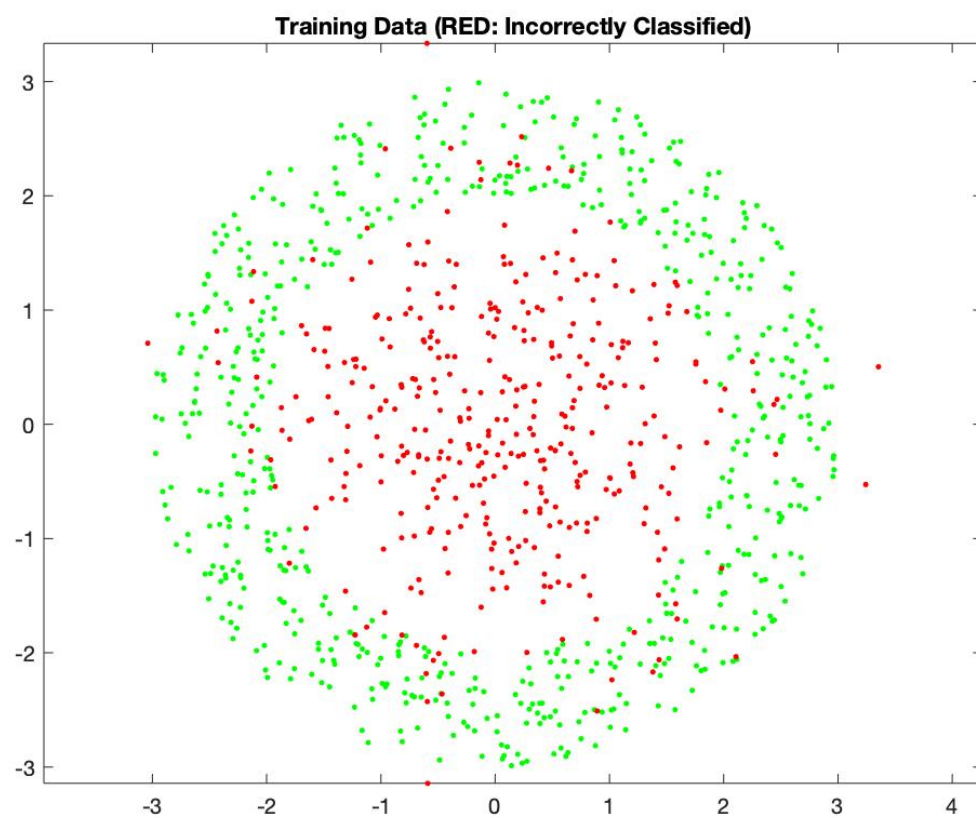


Figure 11: Test Data with Linear SVM




Figure 12: Test data with Gaussian SVM

Our test data results in the following errors:

Linear Test Error: 37.30%

Gaussian Test Error: 5.1%

Once again these errors for the linear match what we would expect based on the priors.

Appendix

1. Problem 1

```
1 %clear all; close all;
2 % Import Image
3 I1 = imread('colorBird.jpg');
4 I2 = imread('colorPlane.jpg');
5 % Get width and height of image
6 Iwidth = length(I1(1,:,1));
7 Iheight = length(I1(:,1,1));
8
9 % Cast image to double for scaling
10 I1 = cast(I1, 'double');
11 Ivec = [reshape(I1(:, :, 1), [], 1), reshape(I1(:, :, 2), [], 1), reshape
        (I1(:, :, 3), [], 1)]; % Reshape image into an nx3 matrix of the
        form [R G B]
12 Ivec = Ivec./256; % Divide by max of uint8 (256) to normalize
        vector
13 indices = find(Ivec(:,1) == Ivec(:,1));
14 % [Vertical distance from corner, Horizontal distance from
        corner]
15 % Vertical and horizontal distance distance starts at 1
16 locFeat = [ceil(indices./(Iwidth)), indices-Iwidth*(ceil(
        indices./(Iwidth))-1)];
17 % Normalize distances
18 locFeat1 = [locFeat(:,1)./max(locFeat(:,1)), locFeat(:,2)./max(
        locFeat(:,2))];
19
20 % Develop Feature Vectors in X as follows [Vertical Distance,
        Horizontal
21 % Distance, R, G, B]
22 X = [locFeat1, Ivec];
23
24 colors(1,:) = [256 0 0];
25 colors(2,:) = [0 256 0];
26 colors(3,:) = [0 0 256];
27 colors(4,:) = [256 256 0];
28 colors(5,:) = [0 256 256];
29 segment = zeros(Iheight, Iwidth, 3);
30 figure(1);
31 for k = 2:5
32     [idx, C] = kmeans(X, k, 'MaxIter', 10000, 'Replicates', 5);
33     idxReshape = reshape(idx, [], Iwidth);
34     segment = zeros(Iheight, Iwidth, 3);
35     for j = 1:k
36         tmp = idxReshape==j;
37         tmp = repmat(tmp, [1 1 3]);
```

```

38         segment(:,:,1) = segment(:,:,1) + colors(j,1).*tmp
           (:,:,1);
39         segment(:,:,2) = segment(:,:,2) + colors(j,2).*tmp
           (:,:,2);
40         segment(:,:,3) = segment(:,:,3) + colors(j,3).*tmp
           (:,:,3);
41     end
42     subplot(2,2,k-1);
43     segment = cast(segment, 'uint8');
44     image(segment);
45     xlabel('X Pixels'), ylabel('Y Pixels'), title(strcat('k-
           means with k = ', num2str(k)));
46 end
47
48 figure(2);
49 for k = 2:5
50     [~, EMPriors, EMmu, EMSigma] = EMforGMM_Edited(length(X
           (:,1)),length(X(1,:)),X',k);
51     idx = evalTopGMM(X, EMPriors, EMmu, EMSigma);
52     idxReshape = reshape(idx,[],Iwidth);
53     segment = zeros(Iheight,Iwidth,3);
54     for j = 1:k
55         tmp = idxReshape==j;
56         tmp = repmat(tmp,[1 1 3]);
57         segment(:,:,1) = segment(:,:,1) + colors(j,1).*tmp
           (:,:,1);
58         segment(:,:,2) = segment(:,:,2) + colors(j,2).*tmp
           (:,:,2);
59         segment(:,:,3) = segment(:,:,3) + colors(j,3).*tmp
           (:,:,3);
60     end
61     subplot(2,2,k-1);
62     segment = cast(segment, 'uint8');
63     image(segment);
64     xlabel('X Pixels'), ylabel('Y Pixels'), title(strcat('k-
           means with k = ', num2str(k)));
65 end
66
67
68 %Begin Plane here
69 Iwidth = length(I2(1,:,1));
70 Iheight = length(I2(:,1,1));
71
72 % Cast image to double for scaling
73 I2 = cast(I2, 'double');
74 Ivec = [reshape(I2(:, :, 1), [], 1), reshape(I2(:, :, 2), [], 1), reshape

```

```

(I2(:,:,3),[],1)]; % Reshape image into an nx3 matrix of the
    form [R G B]
75 Ivec = Ivec./256; % Divide by max of uint8 (256) to normalize
    vector
76 indices = find(Ivec(:,1) == Ivec(:,1));
77 % [Vertical distance from corner, Horizontal distance from
    corner]
78 % Vertical and horizontal distance distance starts at 1
79 locFeat = [ceil(indices./(Iwidth)), indices-Iwidth*(ceil(
    indices./(Iwidth))-1)];
80 % Normalize distances
81 locFeat1 = [locFeat(:,1)./max(locFeat(:,1)), locFeat(:,2)./max(
    locFeat(:,2))];
82
83 % Develop Feature Vectors in X as follows [Vertical Distance,
    Horizontal
84 % Distance, R, G, B]
85 X = [locFeat1, Ivec];
86 segment = zeros(Iheight, Iwidth, 3);
87 figure(3);
88 for k = 2:5
89     [idx, C] = kmeans(X, k, 'MaxIter', 10000);
90     idxReshape = reshape(idx, [], Iwidth);
91     segment = zeros(Iheight, Iwidth, 3);
92     for j = 1:k
93         tmp = idxReshape==j;
94         tmp = repmat(tmp, [1 1 3]);
95         segment(:, :, 1) = segment(:, :, 1) + colors(j, 1).*tmp
            (:, :, 1);
96         segment(:, :, 2) = segment(:, :, 2) + colors(j, 2).*tmp
            (:, :, 2);
97         segment(:, :, 3) = segment(:, :, 3) + colors(j, 3).*tmp
            (:, :, 3);
98     end
99     subplot(2, 2, k-1);
100    segment = cast(segment, 'uint8');
101    image(segment);
102    xlabel('X Pixels'), ylabel('Y Pixels'), title(strcat('k-
        means with k = ', num2str(k)));
103 end
104
105 figure(4);
106 for k = 2:5
107     [~, EMpriors, EMmu, EMSigma] = EMforGMM_Edited(length(X
        (:, 1)), length(X(1, :)), X', k);
108     idx = evalTopGMM(X, EMpriors, EMmu, EMSigma);

```

```

109     idxReshape = reshape(idx,[],Iwidth);
110     segment = zeros(Iheight,Iwidth,3);
111     for j = 1:k
112         tmp = idxReshape==j;
113         tmp = repmat(tmp,[1 1 3]);
114         segment(:,:,1) = segment(:,:,1) + colors(j,1).*tmp
            (:,:,1);
115         segment(:,:,2) = segment(:,:,2) + colors(j,2).*tmp
            (:,:,2);
116         segment(:,:,3) = segment(:,:,3) + colors(j,3).*tmp
            (:,:,3);
117     end
118     subplot(2,2,k-1);
119     segment = cast(segment,'uint8');
120     image(segment);
121     xlabel('X Pixels'), ylabel('Y Pixels'), title(strcat('k-
        means with k = ', num2str(k)));
122 end
123
124
125 function model = evalTopGMM(x, alpha,mu,sigma)
126 tGMM = zeros(size(x,1),length(alpha));
127 for m = 1:length(alpha)
128     tGMM(:,m) = alpha(m)*mvnpdf(x,mu(:,m)',sigma(:,:,m));
129 end
130 [row, col] = find(tGMM == max(tGMM, [], 2));
131 model(row) = col;
132 model = model';
133 % for m = 1:length(alpha)
134 %     tGMM(:,m)
135 % end
136 end
137
138 %%%
139 function gmm = evalGMM(x, alpha,mu,Sigma)
140 gmm = zeros(1,size(x,2));
141 for m = 1:length(alpha) % evaluate the GMM on the grid
142     gmm = gmm + alpha(m)*evalGaussian(x,mu(:,m),Sigma(:,:,m));
143 end
144 end
145 %%%
146 function g = evalGaussian(x,mu,Sigma)
147 % Evaluates the Gaussian pdf N(mu,Sigma) at each column of X
148 [n,N] = size(x);
149 invSigma = inv(Sigma);
150 C = (2*pi)^(-n/2) * det(invSigma)^(1/2);

```

```

151 E = -0.5*sum((x-repmat(mu,1,N)).*(invSigma*(x-repmat(mu,1,N))
    ,1);
152 g = C*exp(E);
153 end

```

2. Problem 2 EM

```

1 function [logLikelihood , alpha ,mu,Sigma] = EMforGMM-Edited(N,d
    , data , parameters)
2 % Generates N samples from a specified GMM
3 % then uses EM algorithm to estimate the parameters
4 % of a GMM that has the same number of components
5 % as the true GMM that generates the samples.
6
7 % close all ,
8 delta = 1e-2; % tolerance for EM stopping criterion
9 regWeight = 5e-3; % regularization parameter for covariance
    estimates
10
11 % Generate samples from a 3-component GMM
12 % alpha_true = [0.2,0.3,0.5];
13 % mu_true = [-10 0 10;0 0 0];
14 % Sigma_true(:,:,1) = [3 1;1 20];
15 % Sigma_true(:,:,2) = [7 1;1 2];
16 % Sigma_true(:,:,3) = [4 1;1 16];
17 % x = randGMM(N,alpha_true ,mu_true ,Sigma_true);
18 % [d,M] = size(mu_true); % determine dimensionality of samples
    and number of GMM components
19 d = d; % Dimension of data coming in.
20 M = parameters;
21 x = data;
22 % Initialize the GMM to randomly selected samples
23 alpha = ones(1,M)/M;
24 shuffledIndices = randperm(N);
25 mu = x(:,shuffledIndices(1:M)); % pick M random samples as
    initial mean estimates
26 [~,assignedCentroidLabels] = min(pdist2(mu',x'),[],1); % assign
    each sample to the nearest mean
27 for m = 1:M % use sample covariances of initial assignments as
    initial covariance estimates
28     Sigma(:, :, m) = cov(x(:, find(assignedCentroidLabels==m)))' +
        regWeight*eye(d,d);
29 end
30 t = 0; %displayProgress(t,x,alpha ,mu,Sigma);
31
32 Converged = 0; % Not converged at the beginning
33 while ~Converged

```

```

34     for l = 1:M % multiply prior with PDF evaluated at certain
        point
35         temp(1,:) = repmat(alpha(1),1,N).*evalGaussian(x,mu(:,l)
            ),Sigma(:, :, l));
36     end
37     plgivenx = temp./sum(temp,1); % Probability of each given
        data
38     alphaNew = mean(plgivenx,2); % New Priors
39     w = plgivenx./repmat(sum(plgivenx,2),1,N); % Probability of
        given data normalized
40     muNew = x*w'; % update means mutliplying data by prior
        distributions
41     for l = 1:M
42         v = x-repmat(muNew(:,l),1,N);
43         u = repmat(w(l,:),d,1).*v;
44         SigmaNew(:, :, l) = u*v' + regWeight*eye(d,d); % adding a
            small regularization term
45     end
46     Dalpha = sum(sum(abs(alphaNew-alpha)));
47     Dmu = sum(sum(abs(muNew-mu)));
48     DSigma = sum(sum(sum(abs(abs(SigmaNew-Sigma)))));
49     alpha = alphaNew; mu = muNew; Sigma = SigmaNew;
50     Converged = ((Dalpha+Dmu+DSigma)<delta); % Check if
        converged
51     t = t+1;
52     % displayProgress(t,x,alpha,mu,Sigma);
53 end
54 logLikelihood = sum(log(evalGMM(x,alpha,mu,Sigma)));
55 %keyboard,
56
57 %%%
58 function displayProgress(t,x,alpha,mu,Sigma)
59 figure(1),
60 if size(x,1)==2
61     subplot(1,2,1), cla,
62     plot(x(1,:),x(2,:), 'b. ');
63     xlabel('x_1'), ylabel('x_2'), title('Data and Estimated GMM
        Contours'),
64     axis equal, hold on;
65     rangex1 = [min(x(1,:)),max(x(1,:))];
66     rangex2 = [min(x(2,:)),max(x(2,:))];
67     [x1Grid,x2Grid,zGMM] = contourGMM(alpha,mu,Sigma,rangex1,
        rangex2);
68     contour(x1Grid,x2Grid,zGMM); axis equal,
69     subplot(1,2,2),
70 end

```



```

71 logLikelihood = sum(log(evalGMM(x, alpha, mu, Sigma)));
72 plot(t, logLikelihood, 'b. '); hold on,
73 xlabel('Iteration Index'), ylabel('Log-Likelihood of Data'),
74 drawnow; pause(0.1),
75
76 %%%
77 function x = randGMM(N, alpha, mu, Sigma)
78 d = size(mu,1); % dimensionality of samples
79 cum_alpha = [0, cumsum(alpha)];
80 u = rand(1,N); x = zeros(d,N); labels = zeros(1,N);
81 for m = 1:length(alpha)
82     ind = find(cum_alpha(m)<u & u<=cum_alpha(m+1));
83     x(:,ind) = randGaussian(length(ind), mu(:,m), Sigma(:, :, m));
84 end
85
86 %%%
87 function x = randGaussian(N, mu, Sigma)
88 % Generates N samples from a Gaussian pdf with mean mu
    covariance Sigma
89 n = length(mu);
90 z = randn(n,N);
91 A = Sigma^(1/2);
92 x = A*z + repmat(mu,1,N);
93
94 %%%
95 function [x1Grid, x2Grid, zGMM] = contourGMM(alpha, mu, Sigma,
    rangex1, rangex2)
96 x1Grid = linspace(floor(rangex1(1)), ceil(rangex1(2)), 101);
97 x2Grid = linspace(floor(rangex2(1)), ceil(rangex2(2)), 91);
98 [h,v] = meshgrid(x1Grid, x2Grid);
99 GMM = evalGMM([h(:)' ; v(:)'], alpha, mu, Sigma);
100 zGMM = reshape(GMM, 91, 101);
101 %figure(1), contour(horizontalGrid, verticalGrid,
    discriminantScoreGrid, [minDSGV
        * [0.9, 0.6, 0.3], 0, [0.3, 0.6, 0.9]*maxDSGV]); % plot equilevel
        contours of the discriminant function
102
103 %%%
104 function gmm = evalGMM(x, alpha, mu, Sigma)
105 gmm = zeros(1, size(x,2));
106 for m = 1:length(alpha) % evaluate the GMM on the grid
107     gmm = gmm + alpha(m)*evalGaussian(x, mu(:,m), Sigma(:, :, m));
108 end
109
110 %%%
111 function g = evalGaussian(x, mu, Sigma)

```

```

112 % Evaluates the Gaussian pdf N(mu,Sigma) at each column of X
113 [n,N] = size(x);
114 invSigma = inv(Sigma);
115 C = (2*pi)^(-n/2) * det(invSigma)^(1/2);
116 E = -0.5*sum((x-repmat(mu,1,N)).*(invSigma*(x-repmat(mu,1,N))),1);
117 g = C*exp(E);

```

3. Problem 2

```

1 % mu(:,1) = [-1;0]; mu(:,2) = [1;0];
2 % Sigma(:,:,1) = [2 0;0 1]; Sigma(:,:,2) = [1 0;0 4];
3 % p = [0.35,0.65]; % class priors for labels 0 and 1
  respectively
4 % % Generate samples
5 % label = rand(1,N) >= p(1); l = 2*(label-0.5);
6 % Nc = [length(find(label==0)),length(find(label==1))]; %
  number of samples from each class
7 % x = zeros(n,N); % reserve space
8 % % Draw samples from each class pdf
9 % for lbl = 0:1
10 %     x(:,label==lbl) = randGaussian(Nc(lbl+1),mu(:,lbl+1),
  Sigma(:,:,lbl+1));
11 % end
12
13 close all, clear all,
14 N=1000; n = 2; K=10;
15 m = 2; % size of feature vector
16 % Class -1 setup
17 mu = zeros(1,m)';
18 Sigma = eye(m);
19
20 % Class +1 setup
21 m = zeros(m,m);
22 m(:,1) = [2,3]';
23 m(:,2) = [-pi, pi]';
24
25 classPriors = [0.35 0.65];
26 classPriors1 = [0.35 0.651];
27 assert(max(cumsum(classPriors)) == 1, 'Priors do not equal 1');
28 thr = [0,cumsum(classPriors1)];
29
30 %Generate Data:
31 % Defined above N = 1000; % Number of samples to generate for
  each set (10,100,1000,10000)
32 u = rand(1,N); L = zeros(1,N); x = zeros(2,N);
33 figure(1), clf, colorList = 'rbgy', hold on;

```

```

34 for l = 1:2
35     indices = find(thr(l)<=u & u<thr(l+1)); % fixed using
        classPriors1 adding a small term to last prior if u
        happens to be precisely 1, that sample will get omitted
        - needs to be fixed
36     L(1,indices) = 1*ones(1,length(indices));
37     if l == 1
38         x(:,indices) = mvnrnd(mu(:,l),Sigma(:, :, l),length(
            indices));
39         plot(x(1,indices),x(2,indices),'.','MarkerFaceColor',
            colorList(l)); axis equal;
40     end
41     if l == 2
42         x(:,indices) = [(m(1,1) + (m(2,1)-m(1,1)).*rand(1,
            length(indices)))', (m(1,2)+ (m(2,2)-m(1,2)).*rand(1,
            length(indices)))']';%r = a + (b-a).*rand(N,1)
43         x(:,indices) = [(x(1,indices).*cos(x(2,indices)))',(x
            (1,indices).*sin(x(2,indices)))']';
44         plot(x(1,indices),x(2,indices),'.','MarkerFaceColor',
            colorList(l)), axis equal;
45         %plot(x(1,indices).*cos(x(2,indices)),x(1,indices).*sin
            (x(2,indices)),'.','MarkerFaceColor',colorList(l));
            axis equal;
46     end
47
48     % axis([-10 10 -10 10]);
49 end
50 xlabel('x1'), ylabel('x2'), legend('Class "-"','Class "+"'),
    title('Data Original Distribution')
51
52 l = 2*(L-1.5);
53
54
55 % Train a Linear kernel SVM with cross-validation
56 % to select hyperparameters that minimize probability
57 % of error (i.e. maximize accuracy; 0-1 loss scenario)
58 dummy = ceil(linspace(0,N,K+1));
59 for k = 1:K, indPartitionLimits(k,:) = [dummy(k)+1,dummy(k+1)];
    end,
60 CList = 10.^linspace(-3,7,11);
61 for CCounter = 1:length(CList)
62     [CCounter,length(CList)],
63     C = CList(CCounter);
64     for k = 1:K
65         indValidate = [indPartitionLimits(k,1):
            indPartitionLimits(k,2)];

```

```

66     xValidate = x(:,indValidate); % Using folk k as
        validation set
67     lValidate = l(indValidate);
68     if k == 1
69         indTrain = [indPartitionLimits(k,2)+1:N];
70     elseif k == K
71         indTrain = [1:indPartitionLimits(k,1)-1];
72     else
73         indTrain = [indPartitionLimits(k-1,2)+1:
            indPartitionLimits(k+1,1)-1];
74     end
75     % using all other folds as training set
76     xTrain = x(:,indTrain); lTrain = l(indTrain);
77     SVMk = fitsvm(xTrain',lTrain,'BoxConstraint',C,'
        KernelFunction','linear');
78     dValidate = SVMk.predict(xValidate'); % Labels of
        validation data using the trained SVM
79     indCORRECT = find(lValidate.*dValidate == 1);
80     Ncorrect(k)=length(indCORRECT);
81     end
82     PCorrect(CCounter)= sum(Ncorrect)/N;
83 end
84 disp(strcat('Minimum error linear CV',num2str(min(Ncorrect))));
85 figure(2), subplot(1,2,1),
86 plot(log10(CList),PCorrect, '.',log10(CList),PCorrect, '-'),
87 xlabel('log_{10} C'),ylabel('K-fold Validation Accuracy
        Estimate'),
88 title('Linear-SVM Cross-Val Accuracy Estimate'), %axis equal,
89 [dummy,indi] = max(PCorrect(:)); [indBestC, indBestSigma] =
        ind2sub(size(PCorrect),indi);
90 CBest= CList(indBestC);
91 SVMBest = fitsvm(x',l', 'BoxConstraint',CBest, 'KernelFunction',
        'linear');
92 d = SVMBest.predict(x'); % Labels of training data using the
        trained SVM
93 indINCORRECT = find(l.*d == -1); % Find training samples that
        are incorrectly classified by the trained SVM
94 indCORRECT = find(l.*d == 1); % Find training samples that are
        correctly classified by the trained SVM
95 figure(2), subplot(1,2,2),
96 plot(x(1,indCORRECT),x(2,indCORRECT),'g.'), hold on,
97 plot(x(1,indINCORRECT),x(2,indINCORRECT),'r.'), axis equal,
98 title('Training Data (RED: Incorrectly Classified)'),
99 disp('Cross-Fold Validation Gaussian Error');
100 pTrainingError = length(indINCORRECT)/N, % Empirical estimate
        of training error probability

```

```

101 Nx = 1001; Ny = 990; xGrid = linspace(-10,10,Nx); yGrid =
    linspace(-10,10,Ny);
102 [h,v] = meshgrid(xGrid,yGrid); dGrid = SVMBest.predict([h(:),v
    (:)]); zGrid = reshape(dGrid,Ny,Nx);
103 figure(2), subplot(1,2,2), contour(xGrid,yGrid,zGrid,0); xlabel
    ('x1'), ylabel('x2'), axis equal,
104 CtrueLinear = CList(indBestC);
105
106 % Train a Gaussian kernel SVM with cross-validation
107 % to select hyperparameters that minimize probability
108 % of error (i.e. maximize accuracy; 0-1 loss scenario)
109 dummy = ceil(linspace(0,N,K+1));
110 for k = 1:K, indPartitionLimits(k,:) = [dummy(k)+1,dummy(k+1)];
    end,
111 CList = 10.^linspace(-1,9,23); sigmaList = 10.^linspace
    (-2,3,23);
112 for sigmaCounter = 1:length(sigmaList)
113     [sigmaCounter,length(sigmaList)],
114     sigma = sigmaList(sigmaCounter);
115     for CCounter = 1:length(CList)
116         C = CList(CCounter);
117         for k = 1:K
118             indValidate = [indPartitionLimits(k,1):
                indPartitionLimits(k,2)];
119             xValidate = x(:,indValidate); % Using folk k as
                validation set
120             lValidate = l(indValidate);
121             if k == 1
122                 indTrain = [indPartitionLimits(k,2)+1:N];
123             elseif k == K
124                 indTrain = [1:indPartitionLimits(k,1)-1];
125             else
126                 indTrain = [indPartitionLimits(k-1,2)+1:
                    indPartitionLimits(k+1,1)-1];
127             end
128             % using all other folds as training set
129             xTrain = x(:,indTrain); lTrain = l(indTrain);
130             SVMk = fitesvm(xTrain',lTrain,'BoxConstraint',C,'
                KernelFunction','gaussian','KernelScale',sigma);
131             dValidate = SVMk.predict(xValidate'); % Labels of
                validation data using the trained SVM
132             indCORRECT = find(lValidate.*dValidate == 1);
133             Ncorrect(k)=length(indCORRECT);
134         end
135         PCorrect(CCounter,sigmaCounter)= sum(Ncorrect)/N;
136     end
end

```

```

137 end
138 disp(strcat('Minimum error Gaussian CV', num2str(min(Ncorrect))
    ));
139
140 figure(3), subplot(1,2,1),
141 contour(log10(CList),log10(sigmaList),PCorrect',50); xlabel('
    log-10 C'), ylabel('log-10 sigma'),
142 title('Gaussian-SVM Cross-Val Accuracy Estimate'), axis equal,
143 [dummy,indi] = max(PCorrect(:)); [indBestC, indBestSigma] =
    ind2sub(size(PCorrect),indi);
144 CBest= CList(indBestC); sigmaBest= sigmaList(indBestSigma);
145 SVMBest = fitcsvm(x',l', 'BoxConstraint',CBest, 'KernelFunction',
    'gaussian', 'KernelScale',sigmaBest);
146 d = SVMBest.predict(x')'; % Labels of training data using the
    trained SVM
147 indINCORRECT = find(1.*d == -1); % Find training samples that
    are incorrectly classified by the trained SVM
148 indCORRECT = find(1.*d == 1); % Find training samples that are
    correctly classified by the trained SVM
149 figure(3), subplot(1,2,2),
150 plot(x(1,indCORRECT),x(2,indCORRECT),'g.'), hold on,
151 plot(x(1,indINCORRECT),x(2,indINCORRECT),'r.'), axis equal,
152 title('Training Data (RED: Incorrectly Classified)'),
153 disp('Cross-Fold Validation Gaussian Error');
154 pTrainingError = length(indINCORRECT)/N, % Empirical estimate
    of training error probability
155 Nx = 10000; Ny = 9900; xGrid = linspace(-10,10,Nx); yGrid =
    linspace(-10,10,Ny);
156 [h,v] = meshgrid(xGrid,yGrid); dGrid = SVMBest.predict([h(:),v
    (:)]); zGrid = reshape(dGrid,Ny,Nx);
157 figure(3), subplot(1,2,2), contour(xGrid,yGrid,zGrid,0); xlabel
    ('x1'), ylabel('x2'), axis equal,
158 CTrue_Gaussian = CList(indBestC);
159 SigmaTrue_Gaussian = sigmaList(indBestSigma);
160
161
162
163 %Generate new Data:
164 % Defined above N = 1000; % Number of samples to generate for
    each set(10,100,1000,10000)
165 u = rand(1,N); L = zeros(1,N); x = zeros(2,N);
166 figure(4),clf, colorList = 'rbgy', hold on;
167 for l = 1:2
168     indices = find(thr(l)<=u & u<thr(l+1)); % fixed using
        classPriors1 adding a small term to last prior if u
        happens to be precisely 1, that sample will get omitted

```

```

169     - needs to be fixed
170     L(1,indices) = l*ones(1,length(indices));
171     if l == 1
172         x(:,indices) = mvnrnd(mu(:,l),Sigma(:, :, l),length(
            indices));
173         plot(x(1,indices),x(2,indices),'.','MarkerFaceColor',
            colorList(l)); axis equal;
174     end
175     if l == 2
176         x(:,indices) = [(m(1,1) + (m(2,1)-m(1,1)).*rand(1,
            length(indices)))', (m(1,2)+ (m(2,2)-m(1,2)).*rand(1,
            length(indices)))']';%r = a + (b-a).*rand(N,1)
177         x(:,indices) = [(x(1,indices).*cos(x(2,indices)))',(x
            (1,indices).*sin(x(2,indices)))']';
178         plot(x(1,indices),x(2,indices),'.','MarkerFaceColor',
            colorList(l)), axis equal;
179         %plot(x(1,indices).*cos(x(2,indices)),x(1,indices).*sin
            (x(2,indices)),'.','MarkerFaceColor',colorList(l));
            axis equal;
180     end
181     % axis([-10 10 -10 10]);
182 end
183 xlabel('x1'), ylabel('x2'), legend('Class "-"','Class "+"'),
    title('Additional Data Original Distribution')
184
185 l = 2*(L-1.5);
186
187 SVMk = fitsvm(xTrain',lTrain','BoxConstraint',CtrueLinear,'
    KernelFunction','linear');
188 testValidate = SVMk.predict(x')'; % Labels of validation data
    using the trained SVM
189 indINCORRECT = find(l.*testValidate == -1); % Find training
    samples that are incorrectly classified by the trained SVM
190 indCORRECT = find(l.*testValidate == 1);
191 Ncorrect(k)=length(indCORRECT);
192 figure(5);
193 plot(x(1,indCORRECT),x(2,indCORRECT),'g. '), hold on,
194 plot(x(1,indINCORRECT),x(2,indINCORRECT),'r. '), axis equal,
195 title('Training Data (RED: Incorrectly Classified)'),
196 disp('Linear SVM Error New Data');
197 pTrainingError = length(indINCORRECT)/N, % Empirical estimate
    of training error probability
198
199
200 SVMk = fitsvm(x',l,'BoxConstraint',CTrue_Gaussian,'

```

```

        'KernelFunction','gaussian','KernelScale',SigmaTrue_Guassian)
    ;
201 dValidate = SVMk.predict(x')'; % Labels of validation data
    using the trained SVM
202 indINCORRECT = find(1.*dValidate == -1); % Find training
    samples that are incorrectly classified by the trained SVM
203 indCORRECT = find(1.*dValidate == 1);
204 Ncorrect(k)=length(indCORRECT);
205 figure(6);
206 plot(x(1,indCORRECT),x(2,indCORRECT),'g.'), hold on,
207 plot(x(1,indINCORRECT),x(2,indINCORRECT),'r.'), axis equal,
208 title('Training Data (RED: Incorrectly Classified)'),
209 disp('Gaussian SVM Error New Data');
210 pTrainingError = length(indINCORRECT)/N, % Empirical estimate
    of training error probability

```